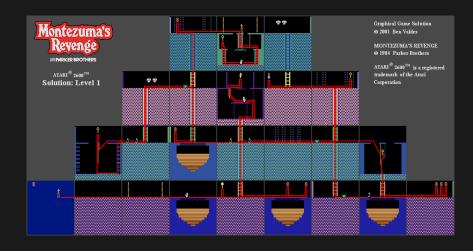


Montezuma's Revenge



Montezuma's Revenge



- very sparse rewards hundrets of steps
- huge state space
- hard exploration
- needs returns back

state of the art score

source: https://paperswithcode.com/sota/ atari-games-on-atari-2600-montezumas-revenge

year	name	score
2015	Deep Reinforcement Learning with Double Q-learning	0
2017	Curiosity-driven Exploration by Self-supervised Prediction ^a	0
2021	MuZero	2500
2018	Count-Based Exploration with Neural Density Models b	3705
2019	Exploration by Random Network Distillation ^c	8152

requires environment state saving/loading

GoExplore* d

2021

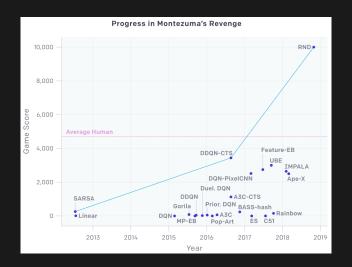
43 000

^ahttps://arxiv.org/abs/1705.05363 ^bhttps://arxiv.org/abs/1703.01310

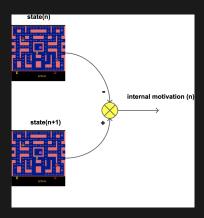
^chttps://arxiv.org/abs/1810.12894

dhttps://arxiv.org/abs/2004.12919

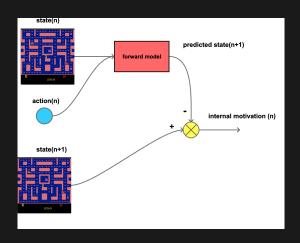
state of the art score



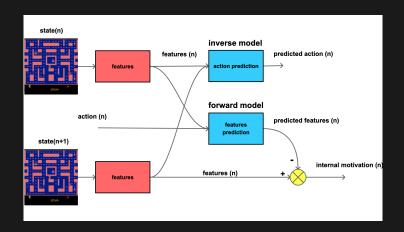
pixel change motivation



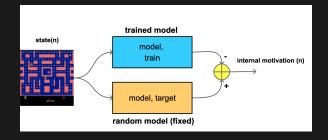
next state prediction



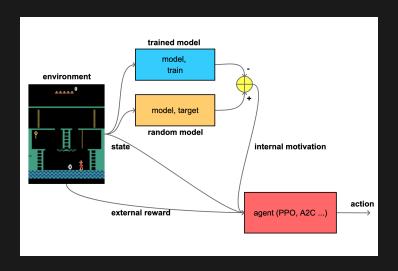
intrinsic curiosity module



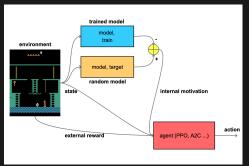
random network distillation



random network distillation

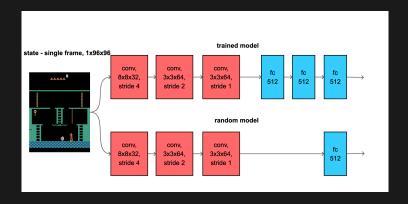


random network distillation

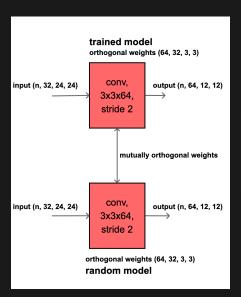


- neural network works as novelty detector
- model learns to imitate random (target) model
- less visited states produce bigger motivation signal
- orthogonal weights initialisation $(g = 2^{0.5})$ for strong signal
- lot of fully connected layers to avoid generalisation
- coupled orthogonal models

random network distillation architecture



coupled RND architecture

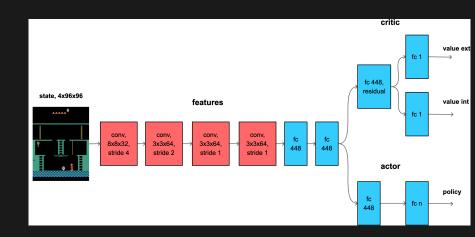


```
def coupled_ortohogonal_init(shape, gain):
    w = torch.zeros((2*shape[0], ) + shape[1:])
    torch.nn.init.orthogonal_(w, gain)

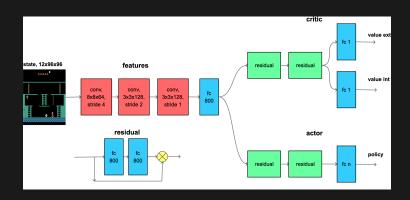
    w = w.reshape((2, ) + shape)
    return w[0], w[1]

wa, wb = coupled_ortohogonal_init((64, 32, 3, 3), 2.0**0.5)
```

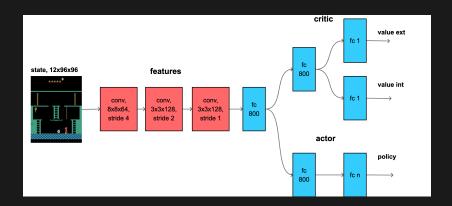
ppo model architecture A



ppo model architecture B

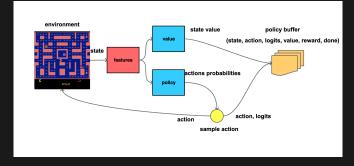


ppo model architecture C

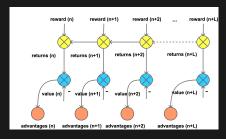


proximal policy optimization - PPO

Schulman, 2017, https://arxiv.org/pdf/1707.06347.pdf

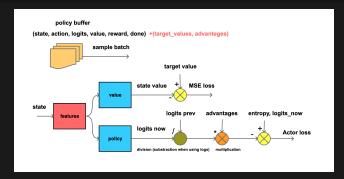


PPO - computing target values and returns



- 1 compute returns (target values), using Q-learning
- 2 compute advantages as difference between returns and current values

PPO - training models



- critic uses common MSE loss
- actors uses loss (excluding clipping terms)

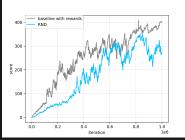
$$\mathcal{L} = rac{1}{N} \sum_{n}^{N} rac{\pi^{now}(a_n|s_n)}{\pi^{prev}(a_n|s_n)} A_n^{\pi^{old}}$$

experiments

- Breakout, without rewards
- Pacman, without rewards
- Montezuma's rewenge (tons of experiments)

breakout

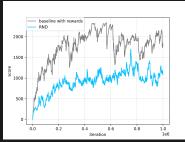




- 1M, 32parallel envs, total 32M steps
- PPO model A
- ext reward weight 2.0
- int reward weight 1.0

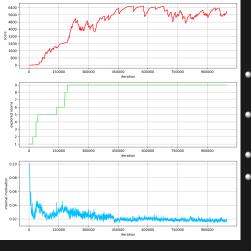
pacman





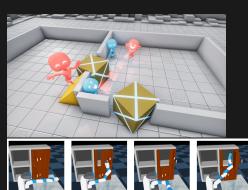
- 1M, 32parallel envs, total 32M steps
- PPO model A
- ext reward weight 2.0
- int reward weight 1.0

results



- 1M steps 20% of original paper
- 128 parallel envs = total128M steps
- score 6400
- 9 rooms explored

Emergent Tool Use From Multi-Agent Autocurricula



- multi-agent robotic environment
- hide and seek
- https: //openai.com/blog/ emergent-tool-use/

https://arxiv.org/abs/1909.07528

Q&A

